

# Sample RAG Document

This is a sample PDF document for testing a Retrieval-Augmented Generation (RAG) pipeline. RAG combines: 1. Document Retrieval 2. Embedding Generation 3. Vector Databases 4. Large Language Models Common tools used in Python RAG systems include LangChain, ChromaDB, Sentence Transformers, and OpenAI APIs. Chunking is important because large documents need to be split into smaller pieces before embeddings are created. Embeddings convert text into vectors so semantic similarity search can be performed efficiently. ChromaDB is a lightweight vector database commonly used for storing embeddings in local RAG applications.